



Susan Li [Follow](#)

Changing the world, one article at a time. Sr. Data Scientist, Toronto Canada. Opinion=my own.

<https://www.linkedin.com/in/susanli/>

Oct 24, 2017 · 8 min read

Linear Regression in Python; Predict The Bay Area's Home Prices



Motivation

In order to predict The Bay area's home prices, I chose the housing price dataset that was sourced from [Bay Area Home Sales Database](#) and [Zillow](#). This dataset was based on the homes sold between January 2013

and December 2015. It has many characteristics of learning, and the dataset can be downloaded from [here](#).

Data Preprocessing

```
import pandas as pd
sf = pd.read_csv('final_data.csv')
sf.head()
```

Unnamed: 0	address	info	z_address	bathrooms	bedrooms	finishedsqft	lastsolddate	lastsoldprice	latitude	longitude	neighborhood	totalrooms
2	Address: 1160 Mission Street #2007 San FranciscoSales price: 1300000Sales date: ...	San FranciscoSales price: 1300000Sales date: ...	1160 Mission St UNIT 2007	2.0	2.0	1043.0	02/17/2016	1300000.0	37.778705	-122.412635	South of Market	4.0
5	Address: 260 King Street #475 San FranciscoSales price: 750000Sales date: 0...	San FranciscoSales price: 750000Sales date: 0...	260 King St UNIT 475	1.0	1.0	903.0	02/17/2016	750000.0	37.777641	-122.393417	South of Market	3.0
7	Address: 560 Missouri Street #B San FranciscoSales price: 1495000Sales date: ...	San FranciscoSales price: 1495000Sales date: ...	560 Missouri St # B	4.0	3.0	1425.0	02/17/2016	1495000.0	37.759198	-122.396516	Potrero Hill	6.0
9	Address: 350 Missouri Street San FranciscoSales price: 2700000Sales date: ...	San FranciscoSales price: 2700000Sales date: ...	350 Missouri St	3.0	3.0	2231.0	02/17/2016	2700000.0	37.761886	-122.396769	Potrero Hill	10.0
11	Address: 3658 Folsom Street San FranciscoSales price: 1530000Sales date: ...	San FranciscoSales price: 1530000Sales date: ...	3658 Folsom St	3.0	3.0	1300.0	02/17/2016	1530000.0	37.740795	-122.413453	Bernal Heights	4.0

There are several features that we do not need, such as “info”, “z_address”, “zipcode”(We have “neighborhood” as a location variable), “zipid” and “zestimate”(This is the price estimated by [Zillow](#), we don’t want our model to be affected by this), so, we will drop them.

```
sf.drop(sf.columns[[0, 2, 3, 15, 17, 18]], axis=1,
        inplace=True)

sf.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 11330 entries, 0 to 11329
Data columns (total 13 columns):
address          11330 non-null object
bathrooms        11330 non-null float64
bedrooms         11330 non-null float64
finishedsqft     11330 non-null float64
lastsolddate     11330 non-null object
lastsoldprice    11330 non-null float64
latitude         11330 non-null float64
longitude        11330 non-null float64
neighborhood     11330 non-null object
totalrooms       11330 non-null float64
usecode          11330 non-null object
yearbuilt        11330 non-null float64
zindexvalue      11330 non-null object
dtypes: float64(8), object(5)
memory usage: 1.1+ MB

```

The data type of “zindexvalue” should be numeric, so let’s change that.

```

sf['zindexvalue'] = sf['zindexvalue'].str.replace(',', '')
sf['zindexvalue'] =
sf['zindexvalue'].convert_objects(convert_numeric=True)

sf.lastsolddate.min(), sf.lastsolddate.max()

```

(‘01/02/2013’, ‘12/31/2015’)

The house sold period in the dataset was between January 2013 and December 2015.

I now use the describe() method to show the summary statistics of the numeric variables.

```
sf.describe()
```

	bathrooms	bedrooms	finishedsqft	lastsoldprice	latitude	longitude	totalrooms	yearbuilt	zindexvalue
count	11330.000000	11330.000000	11330.000000	1.133000e+04	11330.000000	11330.000000	11330.000000	11330.000000	1.133000e+04
mean	1.980229	2.614475	1585.420918	1.263928e+06	37.759711	-122.436518	6.111562	1948.498147	1.320205e+06
std	1.047358	1.299457	921.978245	1.042079e+06	0.025578	0.030743	12.125819	37.911196	5.848170e+05
min	0.500000	0.000000	1.000000	5.350000e+02	37.708170	-122.510726	1.000000	1860.000000	6.881000e+05
25%	1.000000	2.000000	1019.000000	7.292500e+05	37.739286	-122.455157	4.000000	1916.000000	9.829000e+05
50%	2.000000	2.000000	1362.000000	9.900000e+05	37.760513	-122.432510	5.000000	1940.000000	1.211900e+06
75%	2.000000	3.000000	1876.000000	1.450000e+06	37.781386	-122.413359	7.000000	1986.000000	1.480400e+06
max	14.000000	20.000000	27275.000000	2.388900e+07	37.806083	-122.381201	1264.000000	2016.000000	5.333500e+06

The count, mean, min and max rows are self-explanatory. The std shows the standard deviation, and the 25%, 50% and 75% rows show the corresponding percentiles.

To get a feel for the type of data we are dealing with, we plot a histogram for each numeric variable.

```
%matplotlib inline
import matplotlib.pyplot as plt
sf.hist(bins=50, figsize=(20,15))
plt.savefig("attribute_histogram_plots")
plt.show()
```

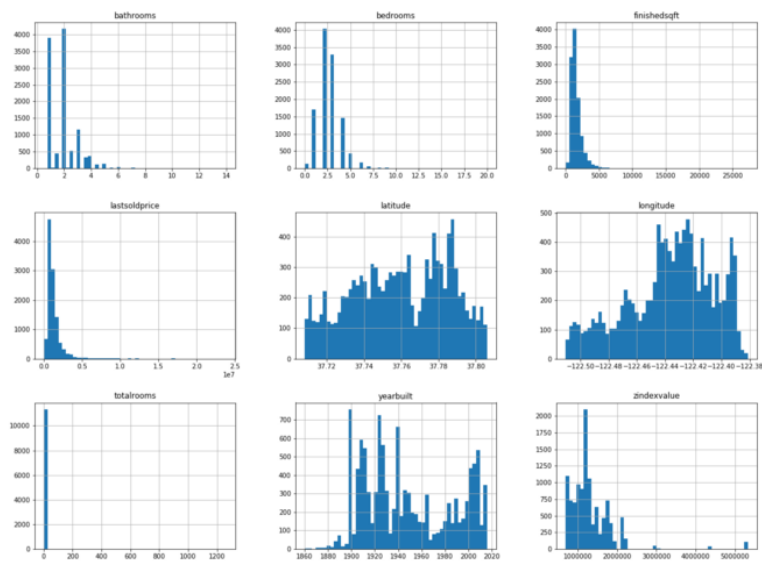


Figure 1. A histogram for each numerical variable

Some of the histograms are a little bit right skewed, but this is not abnormal.

Let's create a scatter plot with latitude and longitude to visualize the data:

```
sf.plot(kind="scatter", x="longitude", y="latitude",  
alpha=0.2)  
plt.savefig('map1.png')
```




Figure 2. A scatter plot of the data

Now let's color code from the most expensive to the least expensive areas:

```
sf.plot(kind="scatter", x="longitude", y="latitude",  
alpha=0.4, figsize=(10,7),  
c="lastsoldprice", cmap=plt.get_cmap("jet"),  
colorbar=True,  
sharex=False)
```




Figure 3. The Bay Area housing prices

This image tells us that the most expensive houses sold were in the north area.

The variable we are going to predict is the “last sold price”. So let’s look at how much each independent variable correlates with this dependent variable.

```
corr_matrix = sf.corr()
corr_matrix["lastsoldprice"].sort_values(ascending=False)
```

```
lastsoldprice    1.000000
finishedsqft     0.647208
bathrooms        0.536880
zindexvalue      0.460429
bedrooms         0.395478
latitude         0.283107
totalrooms       0.093527
longitude        -0.052595
yearbuilt        -0.189055
Name: lastsoldprice, dtype: float64
```

The last sold price tends to increase when the finished sqft and the number of bathrooms go up. You can see a small negative correlation between the year built and the last sold price. And finally, coefficients close to zero indicate that there is no linear correlation.

We are now going to visualize the correlation between variables by using Pandas' `scatter_matrix` function. We will just focus on a few promising variables, that seem the most correlated with the last sold price.

```
from pandas.tools.plotting import scatter_matrix

attributes = ["lastsoldprice", "finishedsqft", "bathrooms",
             "zindexvalue"]
scatter_matrix(sf[attributes], figsize=(12, 8))
plt.savefig('matrix.png')
```




Figure 4. a scatter matrix

The most promising variable for predicting the last sold price is the finished sqft, so let's zoom in on their correlation scatter plot.

```
sf.plot(kind="scatter", x="finishedsqft",
        y="lastsoldprice", alpha=0.5)
plt.savefig('scatter.png')
```





Figure 5. Finished sqft vs. Last Sold Price

The correlation is indeed very strong; you can clearly see the upward trend and that the points are not too dispersed.

Because each house has different square footage and each neighborhood has different home prices, what we really need is the price per sqft. So, we add a new variable “price_per_sqft”. We then check to see how much this new independent variable correlates with the last sold price.

```
sf['price_per_sqft'] =  
sf['lastsoldprice']/sf['finishedsqft']  
corr_matrix = sf.corr()  
corr_matrix["lastsoldprice"].sort_values(ascending=False)
```



```
lastsoldprice    1.000000
finishedsqft     0.647208
bathrooms        0.536880
zindexvalue      0.460429
bedrooms         0.395478
latitude         0.283107
totalrooms       0.093527
price_per_sqft   0.005008
longitude        -0.052595
yearbuilt        -0.189055
Name: lastsoldprice, dtype: float64
```

Unfortunately, the new price_per_sqft variable shows only a very small positive correlation with the last sold price. But we still need this variable for grouping neighborhoods.

There are 71 neighborhoods in the data, and we are going to group them.

```
len(sf['neighborhood'].value_counts())
```

71

The following steps cluster the neighborhood into three groups: 1. low price; 2. high price low frequency; 3. high price high frequency.

buradaki "neighborhood'un
sf'den geleceğini nasıl
algılıyor?

```
freq = sf.groupby('neighborhood').count()['address']
mean = sf.groupby('neighborhood').mean()['price_per_sqft']
cluster = pd.concat([freq, mean], axis=1)
cluster['neighborhood'] = cluster.index

cluster.columns = ['freq', 'price_per_sqft', 'neighborhood']

cluster.describe()
```

	freq	price_per_sqft
count	71.000000	71.000000
mean	159.577465	1664.908308
std	126.572696	3619.277749
min	3.000000	374.201197
25%	67.500000	613.337664
50%	123.000000	756.246284
75%	210.500000	985.156646
max	540.000000	26914.471572

These are the low price neighborhoods:

```
cluster1 = cluster[cluster.price_per_sqft < 756]
cluster1.index
```

*Index(['Bayview', 'Central Richmond', 'Central Sunset', 'Crocker Amazon',
'Daly City', 'Diamond Heights', 'Excelsior', 'Forest Hill',
'Forest Hill Extension', 'Golden Gate Heights', 'Ingleside',
'Ingleside Heights', 'Ingleside Terrace', 'Inner Parkside',
'Inner Richmond', 'Inner Sunset', 'Lakeshore', 'Little Hollywood',
'Merced Heights', 'Mission Terrace', 'Mount Davidson Manor',
'Oceanview', 'Outer Mission', 'Outer Parkside', 'Outer Richmond',
'Outer Sunset', 'Parkside', 'Portola', 'Silver Terrace', 'Sunnyside',
'Visitation Valley', 'West Portal', 'Western Addition',
'Westwood Highlands', 'Westwood Park'],
dtype='object', name='neighborhood')*

These are the high price and low frequency neighborhoods:

```
cluster_temp = cluster[cluster.price_per_sqft >= 756]
cluster2 = cluster_temp[cluster_temp.freq < 123]
cluster2.index
```

*Bayview, - in d, Kaimati, Kullu, Lakewood, NM?,
21, always peak, NM?*

```
Index(['Buena Vista Park', 'Central Waterfront—Dogpatch', 'Corona Heights', 'Haight-Ashbury', 'Lakeside', 'Lone Mountain', 'Midtown Terrace',  
      'North Beach', 'North Waterfront', 'Parnassus—Ashbury', 'Presidio Heights', 'Sea Cliff', 'St. Francis Wood', 'Telegraph Hill', 'Twin Peaks'],  
      dtype='object', name='neighborhood')
```

These are the high price and high frequency neighborhoods:

```
cluster3 = cluster_temp[cluster_temp.freq >=123]  
cluster3.index
```

```
Index(['Bernal Heights', 'Cow Hollow', 'Downtown', 'Eureka Valley—Dolores Heights—Castro', 'Glen Park', 'Hayes Valley', 'Lake', 'Lower Pacific Heights', 'Marina', 'Miraloma Park', 'Mission', 'Nob Hill', 'Noe Valley', 'North Panhandle', 'Pacific Heights', 'Potrero Hill', 'Russian Hill', 'South Beach', 'South of Market', 'Van Ness—Civic Center', 'Yerba Buena'],  
      dtype='object', name='neighborhood')
```

We add a group column based on the clusters:

```
def get_group(x):  
    if x in cluster1.index:  
        return 'low_price'  
    elif x in cluster2.index:  
        return 'high_price_low_freq'  
    else:  
        return 'high_price_high_freq'  
sf['group'] = sf.neighborhood.apply(get_group)
```

After performing the above pre-processing, we do not need the following columns anymore: “address, lastsolddate, latitude, longitude, neighborhood, price_per_sqft”, so, we drop them from our analysis.

```
sf.drop(sf.columns[[0, 4, 6, 7, 8, 13]], axis=1,
        inplace=True)
sf = sf[['bathrooms', 'bedrooms', 'finishedsqft',
        'totalrooms', 'usecode', 'yearbuilt', 'zindexvalue',
        'group', 'lastsoldprice']]
sf.head()
```

	bathrooms	bedrooms	finishedsqft	totalrooms	usecode	yearbuilt	zindexvalue	group	lastsoldprice
0	2.0	2.0	1043.0	4.0	Condominium	2007.0	975700	high_price_high_freq	1300000.0
1	1.0	1.0	903.0	3.0	Condominium	2004.0	975700	high_price_high_freq	750000.0
2	4.0	3.0	1425.0	6.0	Condominium	2003.0	1277600	high_price_high_freq	1495000.0
3	3.0	3.0	2231.0	10.0	SingleFamily	1927.0	1277600	high_price_high_freq	2700000.0
4	3.0	3.0	1300.0	4.0	SingleFamily	1900.0	1248000	high_price_high_freq	1530000.0

Our data looks perfect!

But before we build the model, we need to create dummy variables for these two categorical variables: “usecode” and “group”.

```
X = sf[['bathrooms', 'bedrooms', 'finishedsqft',
        'totalrooms', 'usecode', 'yearbuilt',
        'zindexvalue', 'group']]
Y = sf['lastsoldprice']

n = pd.get_dummies(sf.group)
X = pd.concat([X, n], axis=1)

m = pd.get_dummies(sf.usecode)
X = pd.concat([X, m], axis=1)

drops = ['group', 'usecode']
X.drop(drops, inplace=True, axis=1)

X.head()
```

This is what our data looks like after creating dummy variables:

	bathrooms	bedrooms	finishedsqft	totalrooms	yearbuilt	zindexvalue	high_price_high_freq	high_price_low_freq	low_price	Apartment	Condominium	Coop
0	2.0	2.0	1043.0	4.0	2007.0	975700	1	0	0	0	0	1
1	1.0	1.0	903.0	3.0	2004.0	975700	1	0	0	0	0	1
2	4.0	3.0	1425.0	6.0	2003.0	1277600	1	0	0	0	0	1
3	3.0	3.0	2231.0	10.0	1927.0	1277600	1	0	0	0	0	0
4	3.0	3.0	1300.0	4.0	1900.0	1248000	1	0	0	0	0	0

Train and Build a Linear Regression Model

```
from sklearn.cross_validation import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, Y,
test_size=0.3, random_state=0)

from sklearn.linear_model import LinearRegression
regressor = LinearRegression()
regressor.fit(X_train, y_train)
```

LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)

Done! We now have a working Linear Regression model.

Calculate R squared:

```
y_pred = regressor.predict(X_test)
print('Linear Regression R squared": %.4f' %
regressor.score(X_test, y_test))
```

Linear Regression R squared: 0.5619

So, in our model, 56.19% of the variability in Y can be explained using X. This is not that exciting.

Calculate root-mean-square error (RMSE)

```
import numpy as np
from sklearn.metrics import mean_squared_error
lin_mse = mean_squared_error(y_pred, y_test)
lin_rmse = np.sqrt(lin_mse)
print('Linear Regression RMSE: %.4f' % lin_rmse)
```

Linear Regression RMSE: 616071.5748

Our model was able to predict the value of every house in the test set within \$616071 of the real price.

Calculate mean absolute error (MAE):

```
from sklearn.metrics import mean_absolute_error

lin_mae = mean_absolute_error(y_pred, y_test)
print('Linear Regression MAE: %.4f' % lin_mae)
```

Linear Regression MAE: 363742.1631

Random Forest

Let's try a more complex model to see whether results can be improved —the RandomForestRegressor:

```
from sklearn.ensemble import RandomForestRegressor

forest_reg = RandomForestRegressor(random_state=42)
forest_reg.fit(X_train, y_train)
```

***RandomForestRegressor(bootstrap=True, criterion='mse',
max_depth=None, max_features='auto', max_leaf_nodes=None,
min_impurity_split=1e-07, min_samples_leaf=1,
min_samples_split=2, min_weight_fraction_leaf=0.0,***

n_estimators=10, n_jobs=1, oob_score=False, random_state=42, verbose=0, warm_start=False)

```
print('Random Forest R squared": %.4f' %  
      forest_reg.score(X_test, y_test))
```

Random Forest R squared": 0.6491

```
y_pred = forest_reg.predict(X_test)  
forest_mse = mean_squared_error(y_pred, y_test)  
forest_rmse = np.sqrt(forest_mse)  
print('Random Forest RMSE: %.4f' % forest_rmse)
```

Random Forest RMSE: 551406.0926

Much better! Let's try one more.

Gradient boosting

```
from sklearn import ensemble  
from sklearn.ensemble import GradientBoostingRegressor  
model = ensemble.GradientBoostingRegressor()  
model.fit(X_train, y_train)
```

***GradientBoostingRegressor(alpha=0.9, criterion='friedman_mse',
init=None, learning_rate=0.1, loss='ls', max_depth=3,
max_features=None, max_leaf_nodes=None,
min_impurity_split=1e-07,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=100, presort='auto',
random_state=None, subsample=1.0,
verbose=0, warm_start=False)***


```
print('Gradient Boosting R squared': %.4f' %  
model.score(X_test, y_test))
```

Gradient Boosting R squared”: 0.6616

```
y_pred = model.predict(X_test)  
model_mse = mean_squared_error(y_pred, y_test)  
model_rmse = np.sqrt(model_mse)  
print('Gradient Boosting RMSE: %.4f' % model_rmse)
```

Gradient Boosting RMSE: 541503.7962

These are the best results we have so far, so, I would consider this is our final model.

Feature Importance

We have used 19 features (variables) in our model. Let's find out which features are important and vice versa.

```
feature_labels = np.array(['bathrooms', 'bedrooms',  
    'finishedsqft', 'totalrooms', 'yearbuilt', 'zindexvalue',  
    'high_price_high_freq',  
    'high_price_low_freq', 'low_price', 'Apartment',  
    'Condominium', 'Cooperative',  
    'Duplex', 'Miscellaneous',  
    'Mobile', 'MultiFamily2To4', 'MultiFamily5Plus',  
    'SingleFamily',  
    'Townhouse'])  
importance = model.feature_importances_  
feature_indexes_by_importance = importance.argsort()  
for index in feature_indexes_by_importance:  
    print('{}-{:2f}%'.format(feature_labels[index],  
        (importance[index] *100.0)))
```

```
Apartment-0.00%
MultiFamily5Plus-0.00%
Mobile-0.00%
Miscellaneous-0.00%
Cooperative-0.00%
Townhouse-0.00%
Condominium-0.54%
high_price_low_freq-0.69%
Duplex-1.21%
MultiFamily2To4-1.68%
high_price_high_freq-3.15%
bedrooms-3.17%
low_price-3.57%
SingleFamily-4.12%
yearbuilt-8.84%
totalrooms-9.13%
bathrooms-14.92%
zindexvalue-16.03%
finishedsqft-32.94%
```

The most important features are finished sqft, zindex value, number of bathrooms, total rooms, year built and so on. And the least important feature is Apartment, which means that regardless of whether this unit is an apartment or not, does not matter to the sold price. Overall, most of these 19 features are used.

Your Turn!

Hopefully, this post gives you a good idea of what a machine learning regression project looks like. As you can see, much of the work is in the data wrangling and the preparation steps, and these procedures consume most of the time spent on machine learning.

Now it's time to get out there and start exploring and cleaning your data. Try two or three algorithms, and let me know how it goes.

Source code that created this post can be found [here](#). I would be pleased to receive feedback or questions on any of the above.

